


Kubernetes & Ruby: From Dedicated Servers To Kubernetes

Jakub Kruczek



- General description & ideas (no time to tell about everything, sorry :))
- Join official K8s slack at <http://slack.k8s.io/>
channel **#pl-users** & #kubernetes-ruby

Key presentation agenda

1. Past, Now (& Future)
2. Security (remember!)
3. Let's deploy something
4. Monitoring & Logging
5. Encountered problems

About me

Software Engineer at Boostcom for >3.5y

coding mainly in Ruby, TypeScript and Angular

Started working with Kubernetes >1.5y ago

Deployed, configured and maintain 4 K8s clusters

not counting in deleted ones

I like cats :)

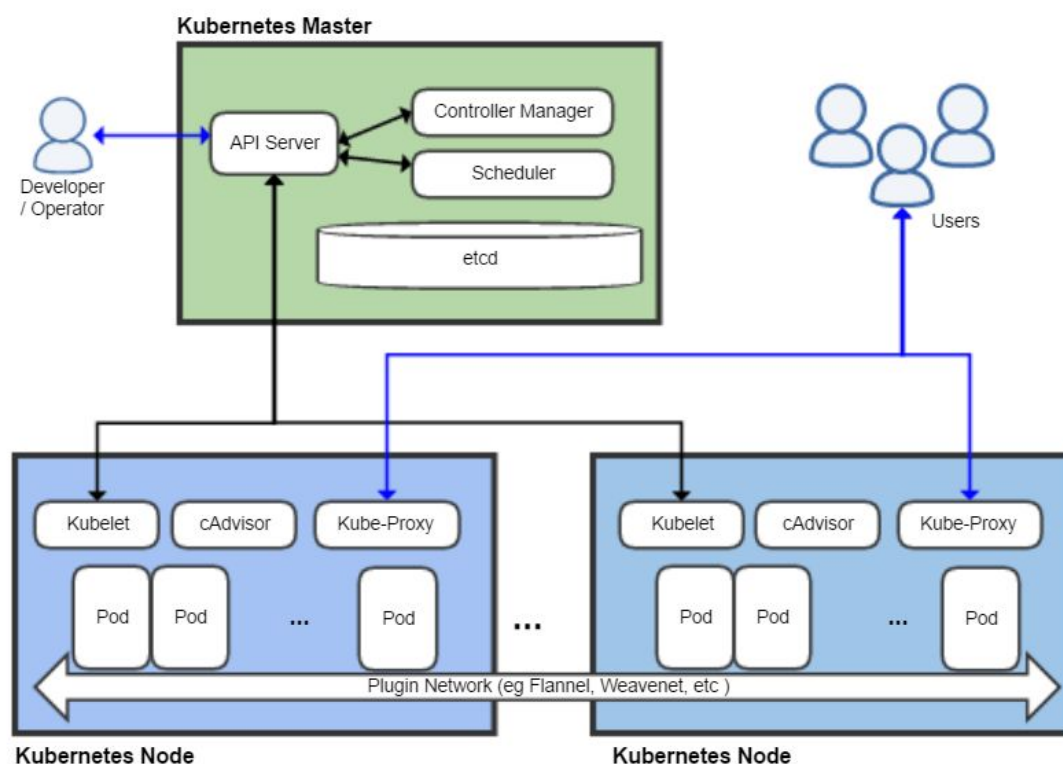


BOOSTCOM



What are the Kubernetes?

Production-Grade Container Orchestration

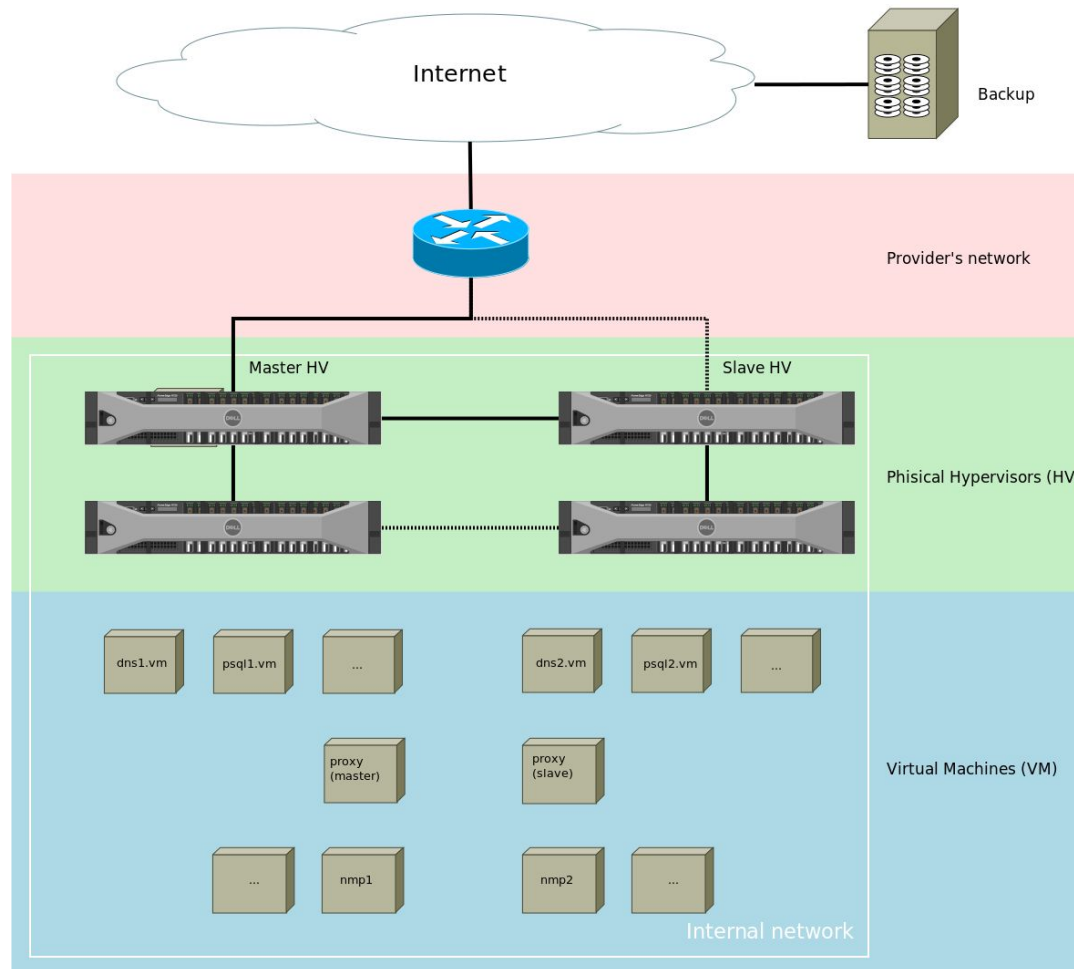


Pod - 1..N container, storage resources, a unique network IP, and additional options

Past, Now (& Future)

Dedicated servers -> Kubernetes

Past - dedicated servers



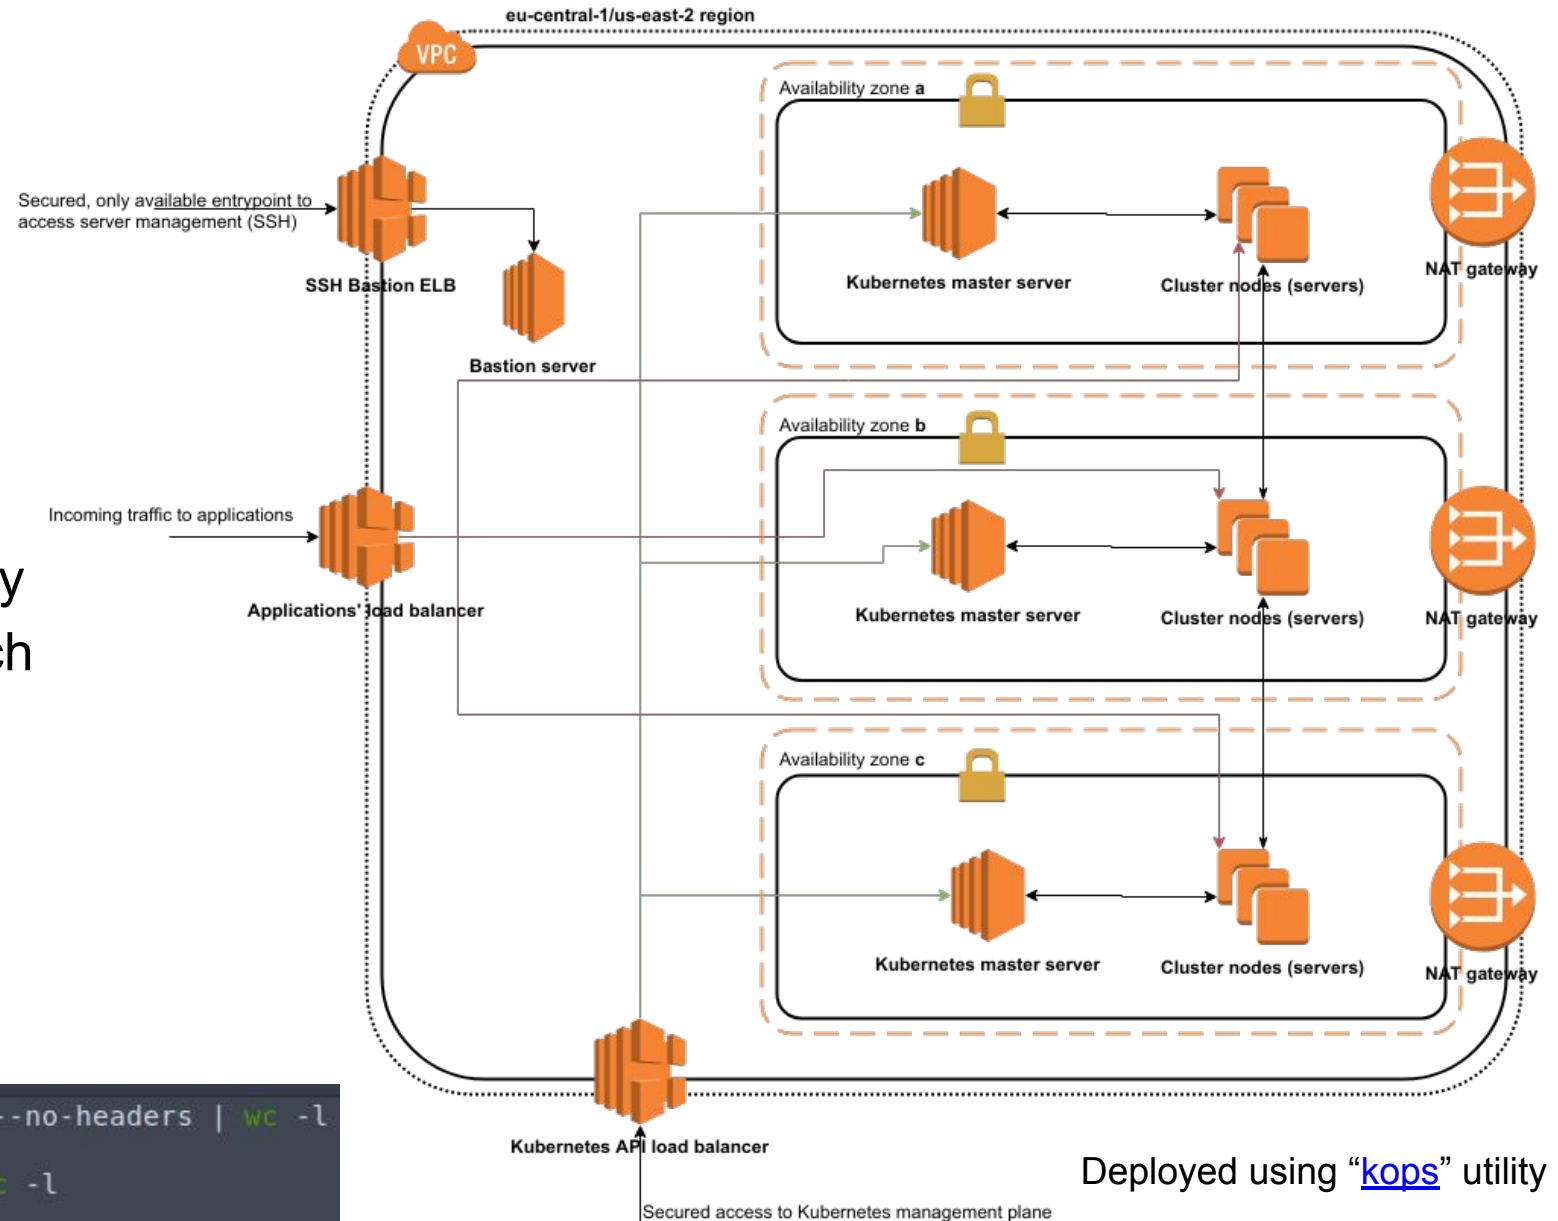
- No real, on-demand scalability
- Single location
- No “Availability Zones”
- Hard to operate and make changes
- GRE tunnels between HV (single tunnel down == everything down...)
- Very custom failover scripts :)
- Very much work needed to make something very simple
- ...

~once a month



Now

- Scalable
- Resilient (and fully HA)
- Multiple locations over the world
- Quite easy to operate by developers without much “servers experience”



Deployed using “[kops](#)” utility

```
[kruczjak:~] % kprod get pods --all-namespaces --no-headers | wc -l
354
[kruczjak:~] % kprod get nodes --no-headers | wc -l
14
[kruczjak:~] %
```

Ruby App migration

Considerations

- Persistent storage (remember about zones (AWS)!)
- Multiple replicas running & HA
- Shutting down & Starting up (scaling, reboot, etc.)
- CPU & RAM usage - Kubernetes requests and limits

Steps

1. Dockerfile!
2. CI/CD scripts
3. Liveness & Readiness probes
4. Kubernetes configuration (yamls, encrypted secrets)
5. ECR repository (just like Dockerhub)
6. Logging adjustments (STDOUT, format)
7. Monitoring!

Total time to migrate single app: ~2 workdays

Ruby App migration - container image

```
Dockerfile x
1 FROM ruby:2.5.1-alpine
2
3 RUN apk --no-cache --update add postgresql-dev tzdata git bind-tools
4
5 # Gemfile for caching
6 ADD Gemfile /app/
7 ADD Gemfile.lock /app/
8 WORKDIR /app
9
10 RUN apk --no-cache --update add --virtual build-dependencies ruby-dev build-base libxml2-dev && \
11     gem install bundler --no-doc && \
12     BUNDLE_FORCE_RUBY_PLATFORM=1 bundle install -j 4 --deployment --without development test && \
13     apk del build-dependencies
14
15 ADD . /app
```

	Alpine	Debian
	398MB	1.3GB
500Mb/s	6.4s	20.8s

Ruby App migration - resources




<https://github.com/linki/chaoskube>

“chaoskube periodically kills random pods in your Kubernetes cluster”

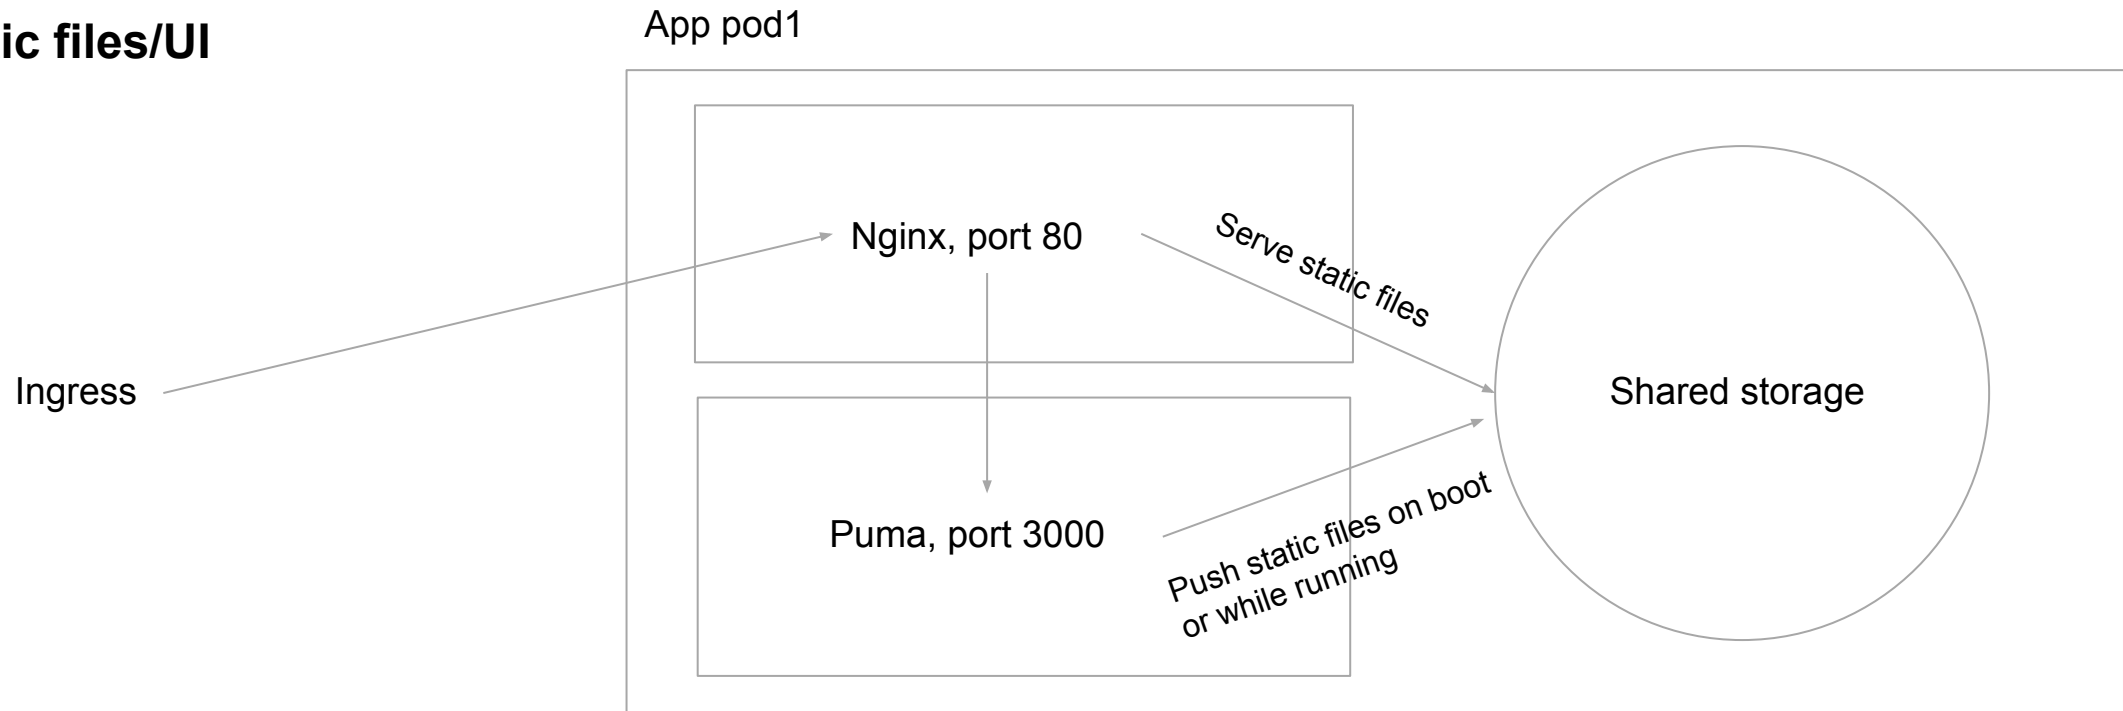
<https://kubernetes.io/docs/concepts/workloads/controllers/jobs-run-to-completion/>

Ruby App migration - static files/UI

API



Static files/UI



Security (remember!)

K8s without proper configuration isn't so secure...

I WANT TO EXPLOIT KUBELET!

Inside pod with PHP application

```
curl -Lk https://<nodeip>:10250/runningpods  
curl -Lk -X POST https://<nodeip>:10250/run/kube-system/<kube-apiserver-podname>/kube-apiserver -d "cmd=cat  
/srv/kubernetes/known_tokens.csv"
```



Solution

- Never allow to connect to node's directly! (hide nodes' IPs)
- Disable "anonymous-auth" on Kubelet
- Enable RBAC authorization on Kubelet

```
/app # curl -k https://172.20.58.165:10250/runningpods/  
Unauthorized/app #
```

AWS & EC2 Roles

```
curl -s http://169.254.169.254/latest/user-data | grep -A 1 channels  
aws s3 ls s3://<s3_name>/
```

Solution

Use e.g. **kube2iam**

```
/app # curl -s http://169.254.169.254/latest/user-data | grep -A 1 channels  
channels:  
- s3://[REDACTED]/addons/bootstrap-channel.yaml  
/app # aws s3 ls s3://[REDACTED]  
  
An error occurred (AccessDenied) when calling the ListObjectsV2 operation: Access Denied  
/app #
```


Ways to access K8s API

- Basic Auth
- Using certificates generated using main CA
- (ServiceAccount) Tokens
- **OpenID Connect Tokens**
 - Google
 - GitHub
 - ...
 - [KeyCloak](#) (2FA, roles, groups)

OAuth2 (e.g. Grafana, Discourse)

<https://github.com/keycloak/keycloak-gatekeeper>

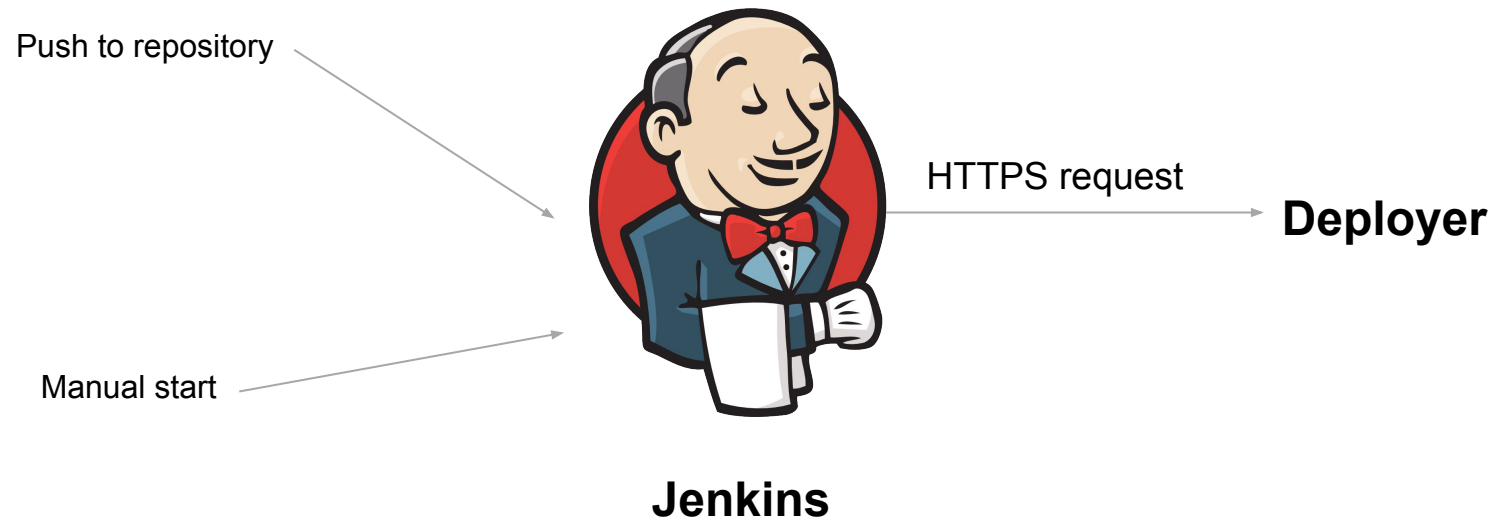
Kibana/ES management dashboard

Prometheus UI (name any app)

Let's deploy something

Manual deployments are soooo funny xD

Deployment overview



- Run tests
- Runs pipeline library:
 - Build Docker image (alpine base)
 - Push image to AWS ECR repository



jenkins APP 12:19

the-big-project-front

Build & push successful 🎉

Configuration:

```
[[awsProfile:boostcom-prod, containerRepoId:492076617898.dkr.ecr.eu-central-1.amazonaws.com], [awsProfile:boostcom-prod, region:us-east-2, containerRepoId:492076617898.dkr.ecr.us-east-2.amazonaws.com]]
```

Branch
master

CommitID
master-547972b74849cb52dc6bf3335ec23e6a506b726e

Build & Deploy

Deployer

```
[kruczzak:~/git/boost/nmp_dashboard/.deploy] develop+* ± tree
├── config.yaml ← shared configuration
├── dev
│   ├── 10-configuration
│   │   ├── a59582db0df375a88d2998291d64fd9efd7335a4464cab58c12993ba2508d85d ← deployment step
│   │   ├── nmp-dashboard-config.configmap.yaml.erb
│   │   ├── nmp-dashboard-nginx-config.configmap.yaml.erb
│   │   ├── nmp-dashboard-secrets.secret.yaml.erb
│   │   └── secrets.ejson
│   ├── 20-predeployment
│   │   └── nmp-dashboard-db-migrate.job.yaml.erb
│   ├── 30-deployment
│   │   ├── nmp-dashboard-puma.ingress.yaml
│   │   ├── nmp-dashboard-puma.web-server-v1.tmpl.yaml.erb
│   │   ├── nmp-dashboard-sidekiq.sidekiq-v1.tmpl.yaml.erb
│   │   └── nmp-dashboard-seeds.job.yaml.erb ← "template"
│   └── 40-postdeployment
│       └── nmp-dashboard-seeds.job.yaml.erb
└── prod
    ├── 10-configuration
    │   ├── nmp-dashboard-config.configmap.yaml.erb
    │   ├── nmp-dashboard-nginx-config.configmap.yaml.erb
    │   ├── nmp-dashboard-secrets.secret.yaml.erb
    │   └── secrets.ejson ← encrypted secrets in app repository
    └── 20-predeployment
        └── nmp-dashboard-db-migrate.job.yaml.erb
```



Deployer APP 15:36

Boostcom Jenkins

reporting_integrations

⚠ Starting deployment

environment

stg

Image

master-c3a82c782fd4ed53187e9a949acfd0263ed69e10

Boostcom Deployer

Boostcom Jenkins

reporting_integrations

Error encountered: **Psych::SyntaxError**

```
/usr/local/lib/ruby/2.4.0/psych.rb:377:in `parse': (<unknown>): could not find expected ':' while scanning a simple key at line 27 column 3
```

environment

stg

Image

master-c3a82c782fd4ed53187e9a949acfd0263ed69e10

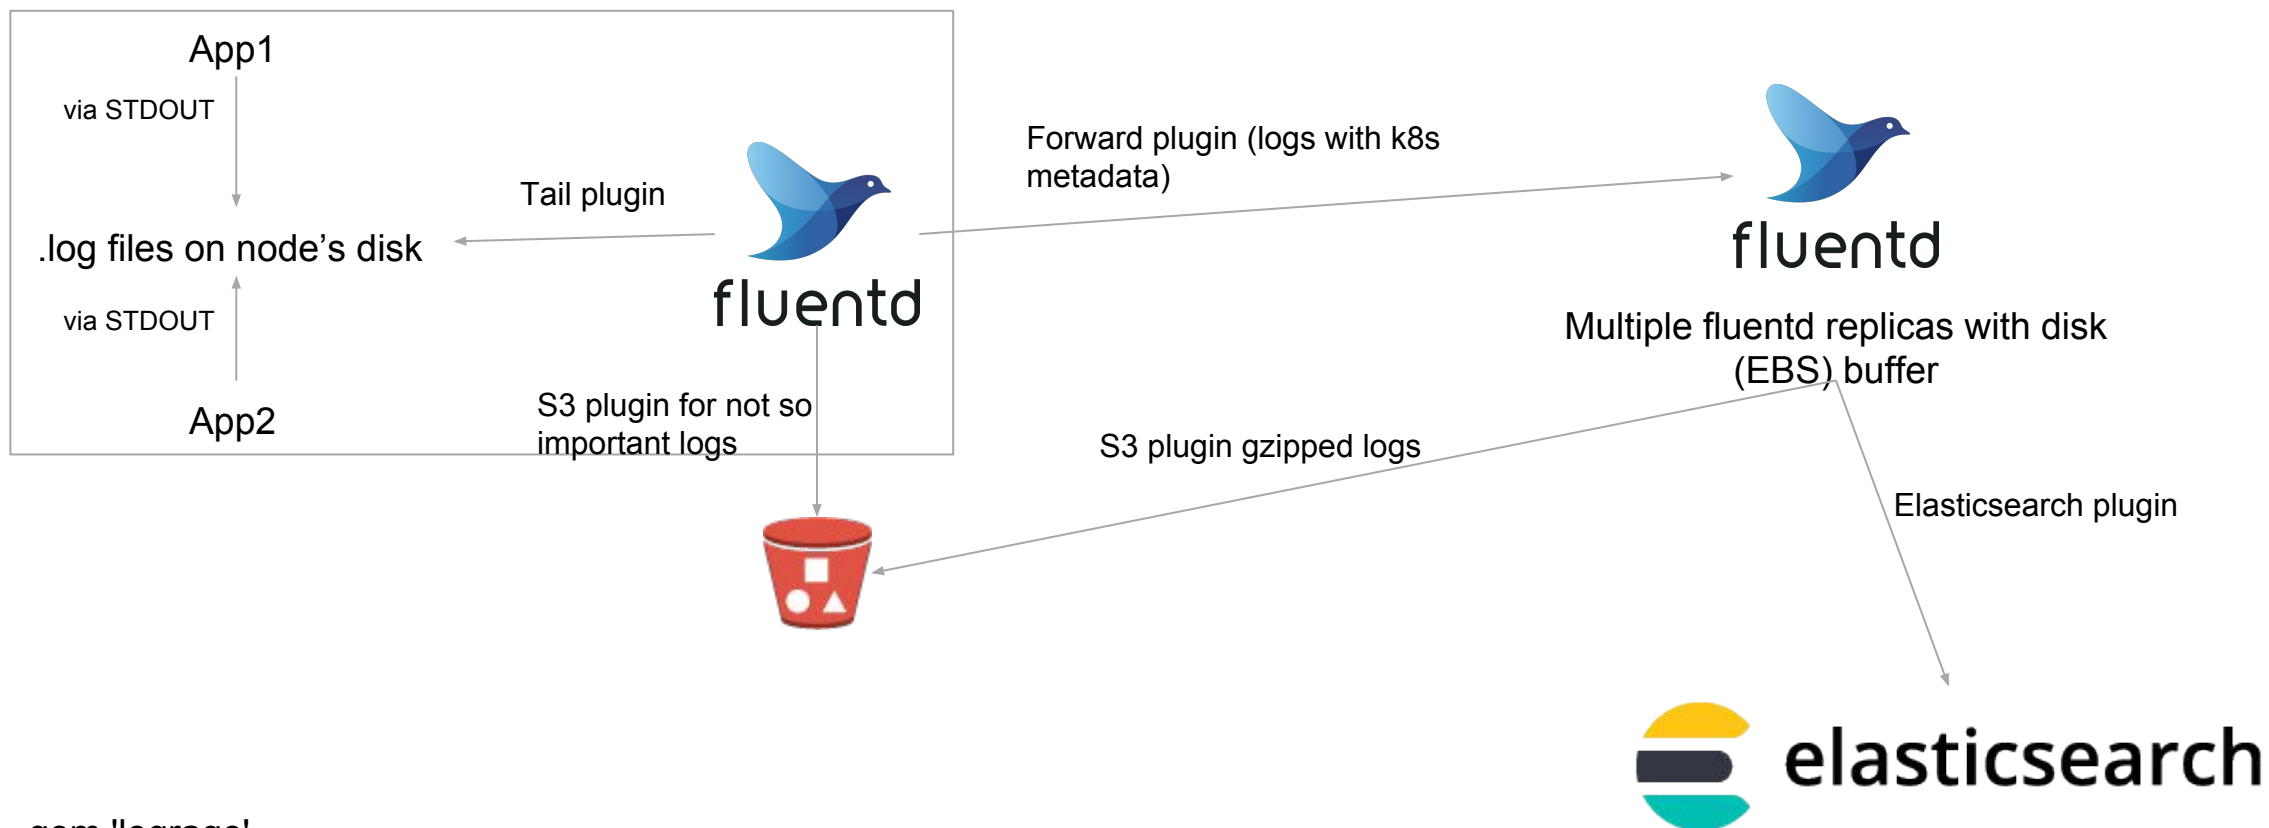
Boostcom Deployer

Monitoring & Logging

Is one thing enough?


































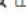











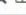





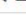










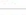




















Logging - EFK stack on k8s

Single node



```
gem 'lograge'  
gem 'logstash-event'  
gem 'logstash-logger'  
gem 'grape_logging' # for grape
```

Logging - EFK stack on k8s

▼ November 18th 2018, 01:32:49   /api/infinity-mall/settings/webform_options 200 infinity-mall [200] GET /api/infinity-mall/settings/webform_options (Api::SettingsController#webform_options)	
Table	JSON View surrounding documents View single document
@timestamp	   * November 18th 2018, 01:32:49.356
@version	   * 1
_id	   * 5b48JGcBHM80Ax0yIypo
_index	   * containers_bl_webforms-api-puma_2018.11.18
_score	   * -
_type	   * fluentd
action	   * webform_options
controller	   * Api::SettingsController
docker.container_id	   * 3f96bcb19c58041e5428782a6eaf216c5620d823fa80cf3a8731002cc4a18577
duration	   * 35.75
format	   * json
host	   * webforms-api-puma-7bddcd5bbd-gqmpz
kubernetes.container_name	   * webforms-api-puma
kubernetes.host	   * ip-172-20-105-25.eu-central-1.compute.internal
kubernetes.labels.pod-template-hash	   * 3688781668
kubernetes.labels.run	   * webforms-api-puma
kubernetes.master_url	   * https://100.64.0.1:443/api
kubernetes.namespace_id	   * fe80faff-0669-11e8-85ca-0a82f5c8558a
kubernetes.namespace_name	   * bl
kubernetes.pod_id	   * 2995c0e6-e865-11e8-9809-0a9195bbab9a
kubernetes.pod_name	   * webforms-api-puma-7bddcd5bbd-gqmpz
log	   * {"method":"GET","path":"/api/infinity-mall/settings/webform_options","format":"json","controller":"Api::SettingsController","action":"webform_options","status":200,"duration":35.75,"view":0.9,"params":{"source":"webforms","slug":"infinity-mall"},"@timestamp":"2018-11-18T00:32:49.356+00:00","@version":"1","message":"[200] GET /api/infinity-mall/settings/webform_options (Api::SettingsController#webform_options)","severity":"INFO","host":"webforms-api-puma-7bddcd5bbd-gqmpz","tags":["149.156.124.9"]}
message	   * [200] GET /api/infinity-mall/settings/webform_options (Api::SettingsController#webform_options)
method	   * GET
params.slug	   * infinity-mall
params.source	   * webforms
path	   * /api/infinity-mall/settings/webform_options

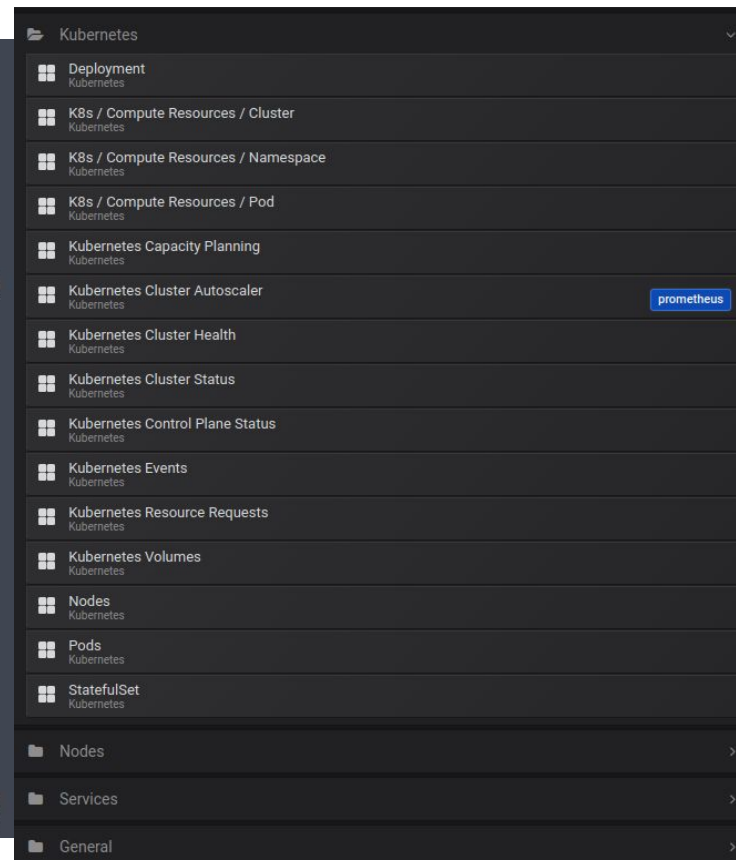
Monitoring - Prometheus & Grafana




<https://github.com/coreos/prometheus-operator> - service discovery, very easy to configure new metrics from pods and alerts

Libs

- https://github.com/discourse/prometheus_exporter
- <https://github.com/yabeda-rb>

```
[kruczjak:~] 130 % kprod get pods -n monitoring
NAME
alertmanager-prometheus-operator-alertmanager-0
alertmanager-prometheus-operator-alertmanager-1
alertmanager-prometheus-operator-alertmanager-2
alerts-keycloak-proxy-7664786d76-5jd49
grafana-6f98fc7c5d-67wq7
kube-slack-7784cdf55c-4x6cz
monitoring-grafana-5bc84c849b-fddjd
postgresql-cluster-production-prometheus-prometheus-postgrc27vb
prometheus-keycloak-proxy-7dbf58645c-4bjq8
prometheus-operator-kube-state-metrics-69995f9794-q4xqr
prometheus-operator-operator-6dc4454b7b-c9n8x
prometheus-operator-prometheus-node-exporter-2p9vz
prometheus-operator-prometheus-node-exporter-775n8
prometheus-operator-prometheus-node-exporter-9dtnw
prometheus-operator-prometheus-node-exporter-crdqq
prometheus-operator-prometheus-node-exporter-czbxp
prometheus-operator-prometheus-node-exporter-dk2dw
prometheus-operator-prometheus-node-exporter-f8z8v
prometheus-operator-prometheus-node-exporter-gthg2
prometheus-operator-prometheus-node-exporter-k8dkv
prometheus-operator-prometheus-node-exporter-kx97x
prometheus-operator-prometheus-node-exporter-pmqxt
prometheus-operator-prometheus-node-exporter-sf5sx
prometheus-operator-prometheus-node-exporter-wtvcv
prometheus-operator-prometheus-node-exporter-zw9v2
prometheus-prometheus-operator-prometheus-0
prometheus-prometheus-operator-prometheus-1
redis-cluster-production-2-prometheus-prometheus-redis-expvs6t6
redis-cluster-production-prometheus-prometheus-redis-expor4nl7l
```




-  **K8s PROD AlertManager** APP 22:56
[FIRING:1] sidekiq-dmp-audience (SidekiqTooManyEnqueuedJobs sidekiq 9292 100.115.240.8:9292 utilities sidekiq-console-web-654b7678-cbttg monitoring/prometheus-operator-prometheus sidekiq-console-metrics-dmp-audience critical)
Sidekiq sidekiq-dmp-audience has queue higher than 300 for 7 minutes. Current: 38376
-  **K8s PROD AlertManager** APP 23:31
[RESOLVED] sidekiq-dmp-audience (SidekiqTooManyEnqueuedJobs sidekiq 9292 100.115.240.8:9292 utilities sidekiq-console-web-654b7678-cbttg monitoring/prometheus-operator-prometheus sidekiq-console-metrics-dmp-audience critical)
Sidekiq sidekiq-dmp-audience has queue higher than 300 for 7 minutes. Current:
-  **K8s PROD AlertManager** APP 12:29
[FIRING:1] node-exporter (metrics monitoring monitoring/prometheus-operator-prometheus prometheus-operator-prometheus-node-exporter)
Load5m on node 172.20.59.52:9100 is high! Currently: 6.4
Load15m on node 172.20.95.178:9100 is high! Currently: 5.02
Load5m on node 172.20.95.178:9100 is high! Currently: 6.14
[RESOLVED] node-exporter (metrics monitoring monitoring/prometheus-operator-prometheus prometheus-operator-prometheus-node-exporter)
Load5m on node 172.20.59.52:9100 is high! Currently: 6.4
Load15m on node 172.20.95.178:9100 is high! Currently: 5.02
Load5m on node 172.20.95.178:9100 is high! Currently: 6.14

Monitoring - Prometheus & Grafana

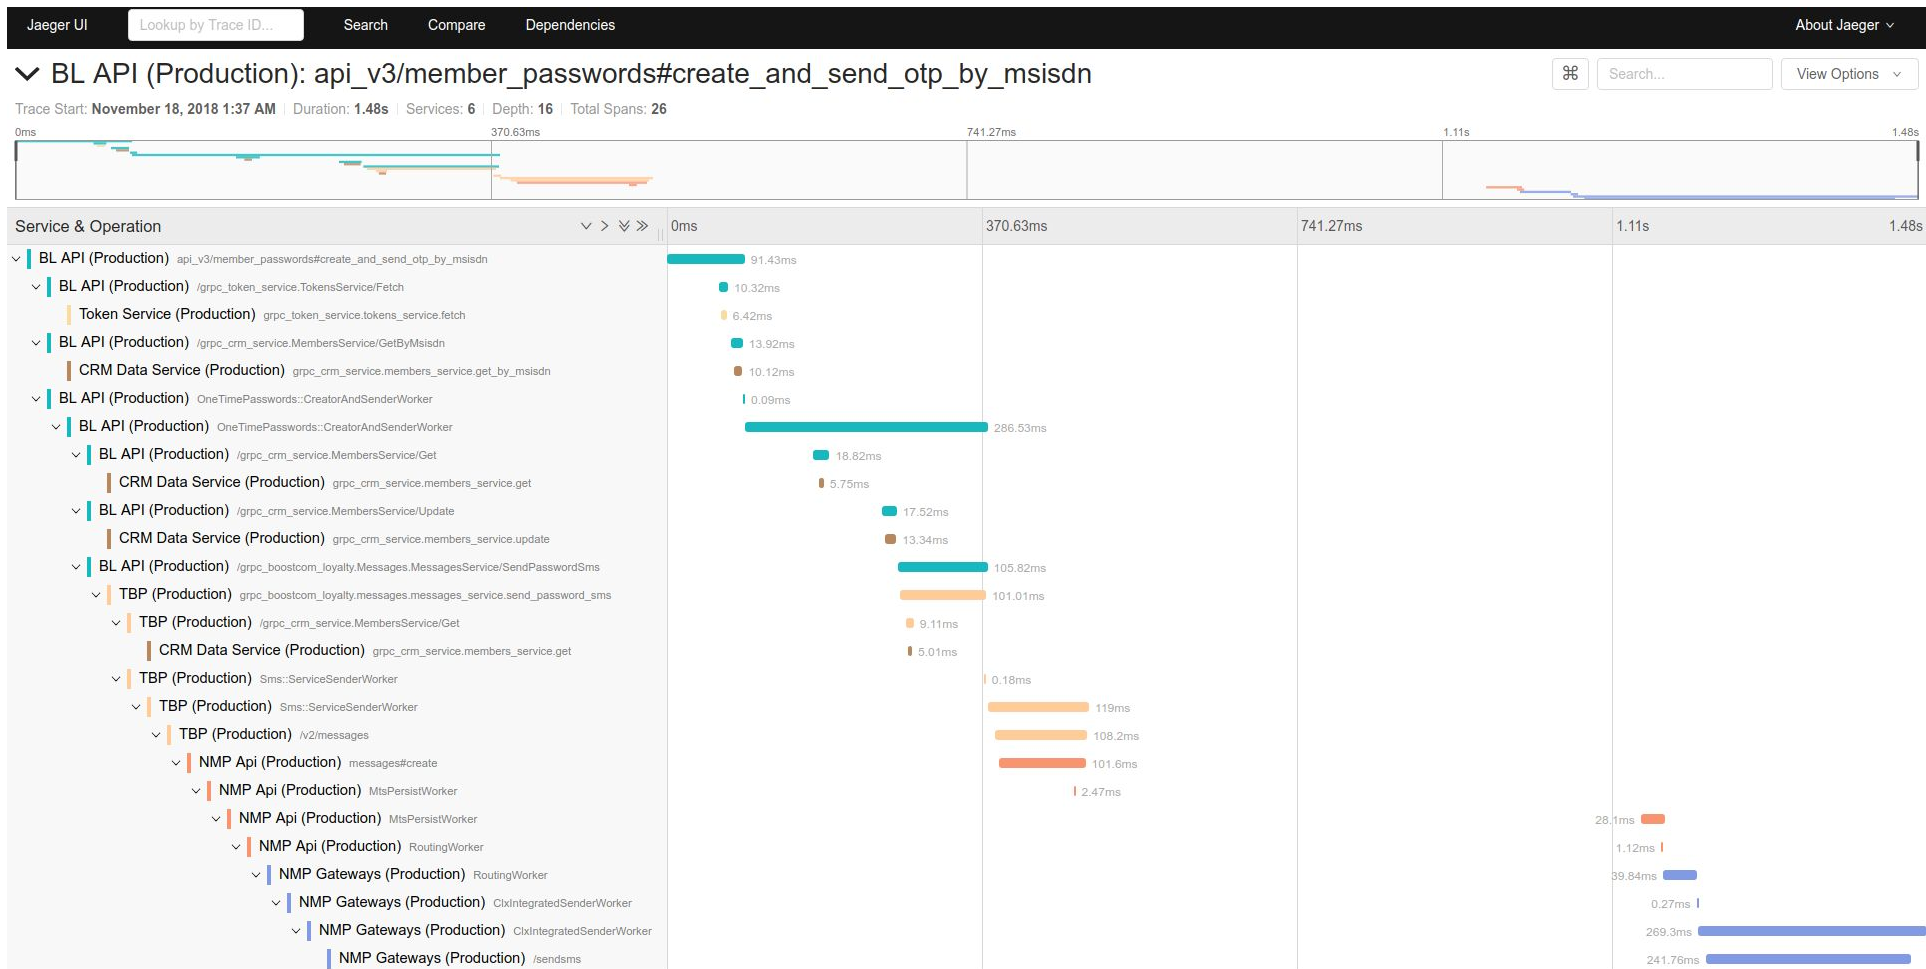


Monitoring - Prometheus & Grafana



Distributed tracing - jaeger

<https://github.com/salemove/jaeger-client-ruby> - example ruby client



Encountered problems

Some things not so great...

Choose “ingress” wisely

nginx-ingress

- C and Golang
- Nginx is not really cloud-first
- Nginx packed inside Golang “configuration wrapper”
- Reconfiguration == nginx reload
- Advanced stuff == manual nginx configuration snippets

```
# ps -eo size,pid,user,command | awk '{ hr=$1/1024 ; printf("%13.6f Mb ",hr) }'
0.000000 Mb COMMAND
0.191406 Mb /usr/bin/dumb-init /nginx-ingress-controller --default-backend=em/nginx-ingress-controller
0.214844 Mb awk { hr=$1/1024 ; printf("%13.6f Mb ",hr) } { for ( x=4 ; x<=
0.328125 Mb /bin/sh
0.328125 Mb /bin/sh
1.074219 Mb ps -eo size,pid,user,command
42.843750 Mb /nginx-ingress-controller --default-backend-service=kube-system/nginx-ingress-controller
55.367188 Mb nginx: master process /usr/sbin/nginx -c /etc/nginx/nginx.conf
318.488281 Mb nginx: worker process is shutting down
318.488281 Mb nginx: worker process is shutting down
318.867188 Mb nginx: worker process is shutting down
319.367188 Mb nginx: worker process
319.367188 Mb nginx: worker process
319.367188 Mb nginx: worker process
319.367188 Mb nginx: worker process
319.367188 Mb nginx: worker process is shutting down
319.367188 Mb nginx: worker process is shutting down
```



traefik

- Golang
- Designed for cloud
- Kubernetes native backend implementation
- Advanced Kubernetes annotations
- Nice UI
- Really easy




The “DNS” bug

Actually - Linux kernel's conntrack bug, while using DNAT and two UDP packets at the same time

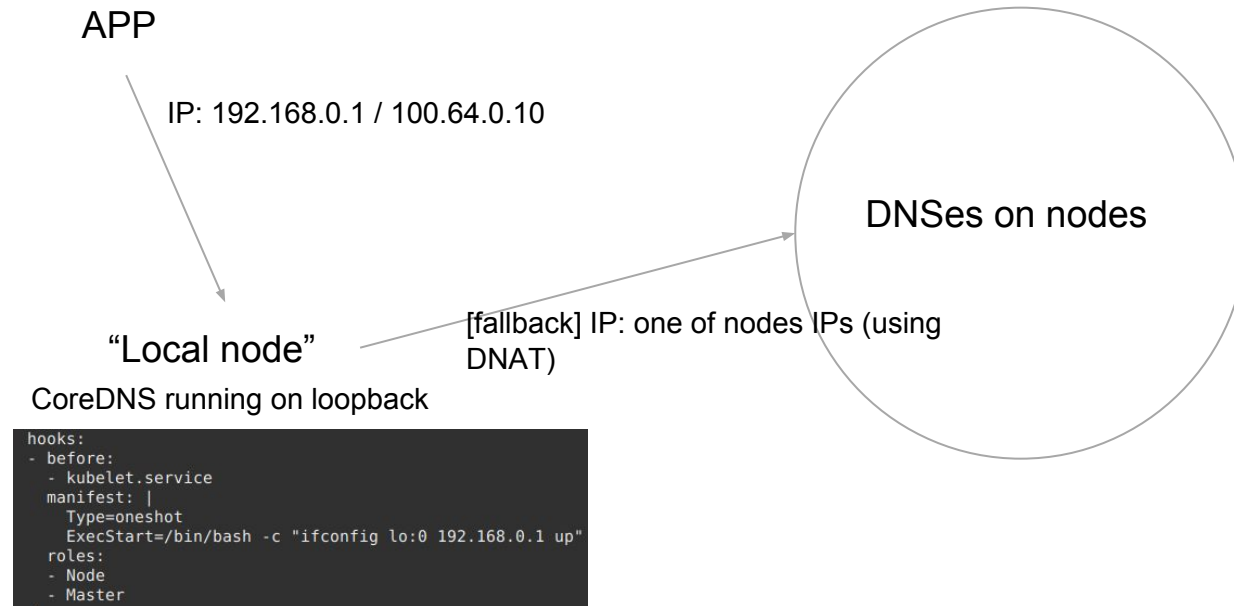
```
irb(main):002:0> 100000.times { time = Time.now; Resolv::DNS.new.getaddress("mp.dev.boost.no"); stoptime = Time.now; p "#{stoptime-time}" if stoptime - time > 0.02 }
"0.055170046"
"5.003717779"
"5.002724427"
"0.069260178"
"0.020280993"
"5.005113935"
"0.052574319"
"0.027052438"
"0.036984"
"5.004995591"
"5.00289057"
"5.006228433"
```

Standard deployment



```
/usr/share/nginx/html # cat /etc/resolv.conf
nameserver 100.64.0.10
search tbp.svc.cluster.local svc.cluster.local cluster.local us-east-2.compute.internal
options ndots:5
```

Workaround deployment



```
hooks:
- before:
  - kubelet.service
  manifest: |
    Type=oneshot
    ExecStart=/bin/bash -c "ifconfig lo:0 192.168.0.1 up"
  roles:
  - Node
  - Master
```

```
/usr/share/nginx/html # cat /etc/resolv.conf
nameserver 192.168.0.1
nameserver 100.64.0.10
search tbp.svc.cluster.local svc.cluster.local cluster.local eu-central-1.compute.internal
options ndots:5
```

Summary

Summary

Pros

- Simple to operate and configure
- Super scalable
- Server downtimes reduced almost to 0
- Very popular with big community and mass of dedicated services
- Versatile
- Simple app deployments (compared to dedicated servers)

Cons


- You have to learn basics first (which can be a bit hard) [developers]
- You have to know something about networking. [maintaining/operating]
- Not for single app (overhead)
- Not so easy to deploy on barebone servers (with good configuration)

We're hiring!

Know Ruby/PHP/Angular?

Send your CV to techjobs@boostcom.no



 +47 73 60 60 23

 contact@boostcom.no

 boostcom.no